

PART A: TIME ANALYSIS

A.1: Number of transactions per month

To get the monthly transaction total, we used the mapper's 'datetime.utcfromtimestamp' functionality to extract the date in the desired format, which includes month and year.

In the reducer section, the counts for the various keys are then added together.

Execution command: `python no_trans.py -r hadoop --output-dir trans_per_month --no-cat-output hdfs://andromeda.eecs.qmul.ac.uk/data/ethereum/transactions/part-*.csv`

```
jmp01@it1211 ~/ecs765/ass1/partA> cat no_trans.py
from mrjob.job import MRJob
from datetime import datetime
import time
class no_trans(MRJob):

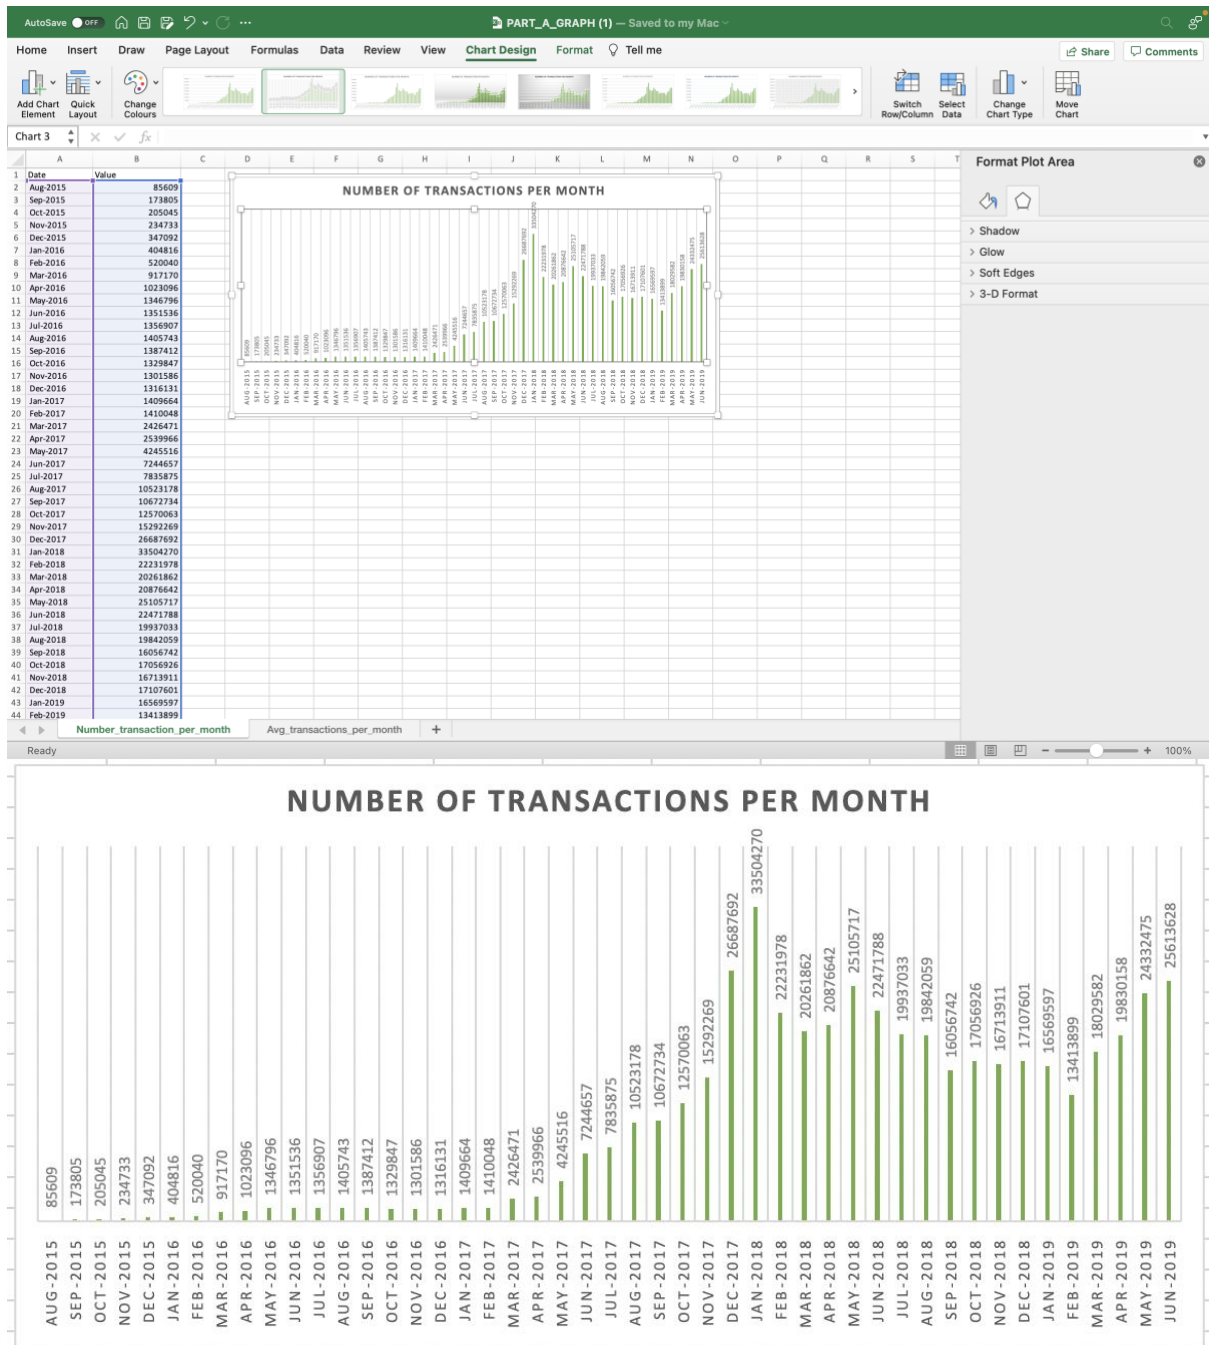
    def mapper(self, _, line):
        try:
            fields = line.split(',')
            if len(fields) == 7:
                day = (datetime.utcfromtimestamp(int(fields[6])).strftime('%Y-%m'))
                value = int(fields[3])
                yield (day,1)
        except :
            pass
    def reducer(self,u,t):
        yield u, sum(t)

if __name__ == '__main__':
    no_trans.run()

# python no_trans.py -r hadoop --output-dir trans_per_month --no-cat-output hdfs://andromeda.eecs.qmul.ac.uk/data/ethereum/transactions/part-*.csv

# job: http://andromeda.student.eecs.qmul.ac.uk:8088/proxy/application_1648683650522_6967/
Job Job_1648683650522_6967 completed successfully
Output directory: hdfs:///user/jmp01/trans_per_month
Counters: 55
File Input Format Counters
  Bytes Read=65309532836
File Output Format Counters
  Bytes Written=860
File System Counters
  FILE: Number of bytes read=321522456
  FILE: Number of bytes written=888562378
  FILE: Number of large read operations=0
  FILE: Number of read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=65309715788
  HDFS: Number of bytes read erasure-coded=0
  HDFS: Number of bytes written=860
  HDFS: Number of large read operations=0
  HDFS: Number of read operations=3277
  HDFS: Number of write operations=4
Job Counters
  Data-local map tasks=986
  Launched map tasks=1089
  Launched reduce tasks=2
  Rack-local map tasks=103
  Total megabyte-milliseconds taken by all map tasks=63529036800
  Total megabyte-milliseconds taken by all reduce tasks=12193628160
  Total time spent by all map tasks (ms)=12480015
  Total time spent by all maps in occupied slots (ms)=62040075
  Total time spent by all reduce tasks (ms)=2381568
  Total time spent by all reduces in occupied slots (ms)=11907840
  Total vcore-milliseconds taken by all map tasks=12480015
  Total vcore-milliseconds taken by all reduce tasks=2381568
Map-Reduce Framework
  CPU time spent (ms)=6358430
  Combine input records=0
  Combine output records=0
  Failed Shuffles=0
  GC time elapsed (ms)=153877
  Input split bytes=182952
  Map input records=486522454
  Map output bytes=5838256380
  Map output materialized bytes=321556430
  Map output records=486521365
  Merged Map outputs=2178
  Peak Map Physical memory (bytes)=668643328
  Peak Map Virtual memory (bytes)=2965000192
  Peak Reduce Physical memory (bytes)=1137086464
  Peak Reduce Virtual memory (bytes)=3019788288
  Physical memory (bytes) snapshot=675127648256
  Reduce input groups=47
  Reduce input records=486521365
  Reduce output records=47
  Reduce shuffle bytes=321556430
  Shuffled Maps =2178
  Spilled Records=973042730
  Total committed heap usage (bytes)=694700690432
  Virtual memory (bytes) snapshot=2982787223552
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
```

Bargraph:



A.2: Average value of transactions per month

In the mapper, we acquired the relevant 'value' field from the transaction's schema along with the date and cumulated these in the reducer phase, like the prior code.

Execution command: `python trans_value.py -r hadoop --output-dir each_month_transaction -no-cat-output hdfs://andromeda.eecs.qmul.ac.uk/data/ethereum/transactions/part-*.csv`

```
[jmp01@itl211 ~/ecs765/ass1/partA> cat trans_value.py
from mrjob.job import MRJob
from mrjob.step import MRStep
import fileinput
import sys
from time import gmtime, strftime, struct_time
from datetime import datetime

class trans_value(MRJob):

    def mapper(self, _, line):
        try:
            fields = line.split(',')
            if len(fields)==7:
                timestamp_val=(datetime.utcfromtimestamp(int(fields[6])).strftime("%b,%Y "))
                value=int(fields[3])
                yield(timestamp_val,value)

        except:
            pass

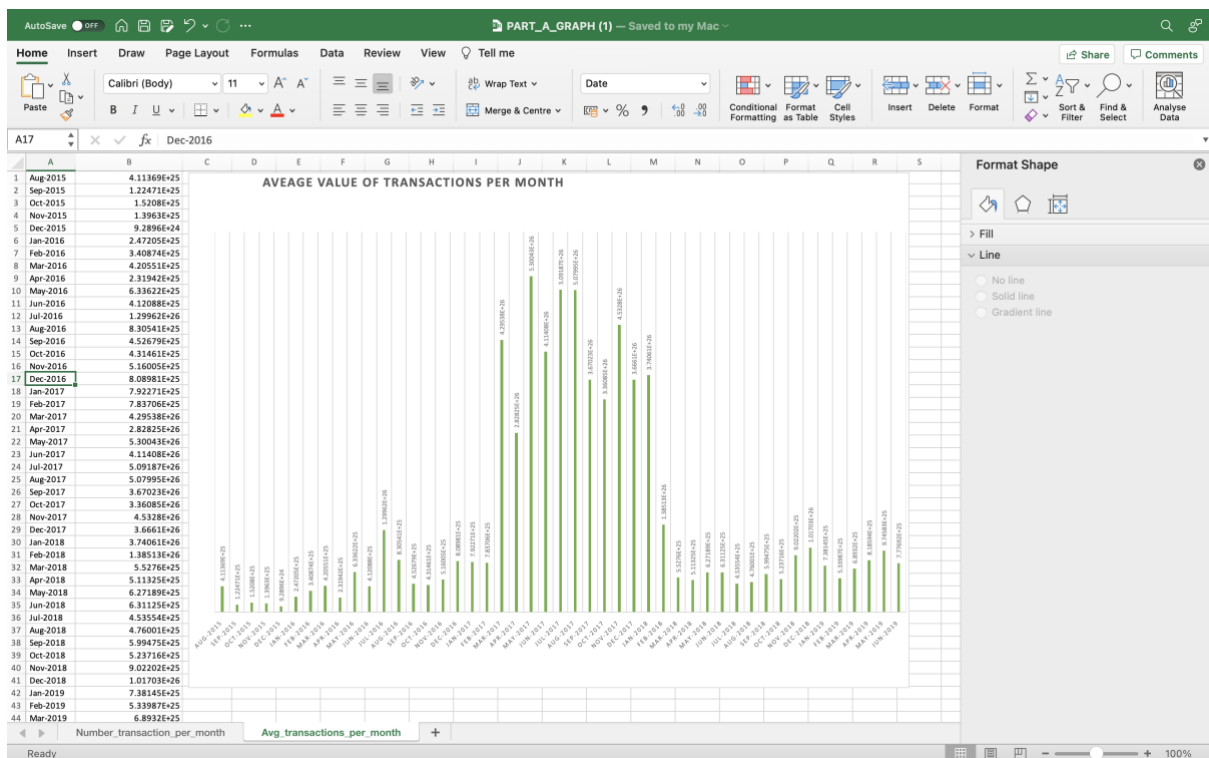
    def combiner(self,timestamp_val,value):
        yield(timestamp_val,sum(value))


    def reducer(self,timestamp_val,value):
        yield(timestamp_val, sum(value))

if __name__ == "__main__":

    trans_value.run()

# python trans_value.py -r hadoop --output-dir each_month_transaction --no-cat-output hdfs://andromeda.eecs.qmul.ac.uk/data/ethereum/transactions/part-*.csv
# job: http://andromeda.student.eecs.qmul.ac.uk:8088/proxy/application_1648683650522_7813/
```





PART B: TOP TEN MOST POPULAR SERVICES

1. Aggregation

The input transaction file is read and split at ',' in the first step of aggregation. We also get the key: to address and the value: value. The entire value for each unique key (to address) is then added.

Execution Command: `python agg_trans.py -r hadoop --output-dir agg_trans --no-cat-output hdfs://andromeda.eecs.qmul.ac.uk/data/ethereum/transactions/part-*.csv`

```

jmp01@itl211 ~/ecs765/ass1/partB> cat agg_trans.py
from mrjob.job import MRJob
from mrjob.step import MRStep
import fileinput
import sys
from time import gmtime, strftime, struct_time
from datetime import datetime

class trans_per_m(MRJob):

    def mapper(self, _, line):
        try:
            fields = line.split(',')
            if len(fields)==7:
                to_address=fields[2]
                value=int(fields[3])
                yield(to_address,value)

        except:
            pass

    def reducer(self,to_address,value):
        yield(to_address, sum(value))

if __name__ == "__main__":

    trans_per_m.run()

# python agg_trans.py -r hadoop --output-dir agg_trans --no-cat-output hdfs://andromeda.eecs.qmul.ac.uk/data/ethereum/transactions/part-*.csv

```

2. Repartition join between aggregate and contracts

'/user/jmp01/aggregate transactions/' and 'data/ethereum/contracts/' are the two input source directories. We look to split the input in the mapper phase based on whether the delimiter is 't' (in aggregate transactions with 2 fields) OR ',' (in contracts with 5 fields). We also get key(address) and values (block number).

Only when a key(address) receives values from both 'aggregate transactions' and 'contracts' will the reducer yield in the subsequent reducer phase (address, value). As a result, only smart contracts are allowed to pass.

The result includes a list of smart contracts as well as their total values.

Execution command: `python join_task2.py -r hadoop --output-dir part_b_joined_records --no-cat-output`
`hdfs://andromeda.eecs.qmul.ac.uk/data/ethereum/contracts/`
`hdfs://andromeda.eecs.qmul.ac.uk/user/jmp01/aggregate_transactions/`

```
jmp01@itl211 ~/ecs765/ass1/partB> cat join_task2.py
from mrjob.job import MRJob
from mrjob.step import MRStep
import fileinput
import sys
from time import gmtime, strftime, struct_time
from datetime import datetime

class transaction_per_month(MRJob):
    def mapper(self, _, line):
        try:
            if (len(line.split('\t'))==2):
                fields = line.split('\t')
                join_key = fields[0]
                join_val = int(fields[1])
                yield (join_key, (join_val, 1))

            elif (len(line.split(','))==5):
                fields = line.split(',')
                join_key = fields[0]
                join_val = fields[3]
                yield (join_key, (join_val, 2))

        except:
            pass

    def reducer(self, address, values):
        contracts = []
        aggregate_file = None

        for value in values:
            if value[1]==1:
                aggregate_file=value[0]
            elif value[1]==2:
                contracts.append(value[0])
            if aggregate_file != None and len(contracts)!=0:
                yield (address, aggregate_file)

if __name__ == "__main__":
    transaction_per_month.run()

# python join_task2.py -r hadoop --output-dir joined_records --no-cat-output hdfs://andromeda.eecs.qmul.ac.uk/data/ethereum/contracts/ hdfs://andromeda.eecs.qmul.ac.uk/user/kg003/aggregate_transactions/
```

3. Finding top 10

If the input file is divided along 't' and has two fields, the mapper reads it line by line from the preceding job (/user/jmp01/part b joined records/). We get value(None) and key(None) (address, total transaction value).

All pairs are sorted in descending order in the reducer phase using the sort function. Finally, we run through the sorted data 10 times before stopping.

Execution Command: `python part_B_top10.py -r hadoop --output-dir part_B_top10 --no-cat-output`
`hdfs://andromeda.eecs.qmul.ac.uk/user/jmp01/part_b_joined_records`

```
jmp01@itl1211 ~/ecs765/ass1/part8> cat part_B_top10.py
from mrjob.job import MRJob

class PartB(MRJob):

    def mapper(self,_,line):
        try:
            fields = line.split('\t')
            if len(fields) == 2:
                add = fields[0]
                agg = float(fields[1])
                yield (None,(add, agg))
        except:
            pass

    def combiner(self,key,val):
        sorted_values = sorted(val,reverse=True, key=lambda tup:tup[1])

        i=0
        for v in sorted_values:
            yield ("Top",v)
            i += 1
            if i >= 10:
                break

    def reducer(self,key,val):
        sorted_values = sorted(val, reverse=True, key=lambda tup: tup[1])

        i = 0
        for v in sorted_values:
            yield ((v[0], v[1]))
            i+=1
            if i>=10:
                break

if __name__=='__main__':
    PartB.run()

#python part_B_top10.py -r hadoop --output-dir part_B_top10 --no-cat-output hdfs://andromeda.eecs.qmul.ac.uk/user/jmp01/part_b_joined_records
```

Output:

```
0xaa1a6e3e6ef20068f7f8d8c835d2d22fd5116444 8.415510080996593e+25
0x3f5ce5fbfe3e9af3971dd833d26ba9b5c936f0be 5.834331302252918e+25
0x32be343b94f860124dc4fee278fdecbd38c102d88 5.432072061064933e+25
0xfa52274dd61e1643d2205169732f29114bc240b3 4.578748448318936e+25
0x7727e5113d1d161373623e5f49fd568b4f543a9e 4.56206240013507e+25
0x209c4784ab1e8183cf58ca33cb740efbf3fc18ef 4.317035609226244e+25
0x876eabf441b2ee5b5b0554fd502a8e0600950cfa 4.015767887861935e+25
0x6fc82a5fe25a5cdb58bc74600a40a69c065263f8 2.706892158201953e+25
0x2910543af39aba0cd09dbb2d50200b3e800a63d2 2.6386022236908247e+25
0xcafb10ee663f465f9d10588ac44ed20ed608c11e 2.3078610109547e+25
```

PARTC: TOP TEN MOST ACTIVE MINERS

The mapper phase reads the file line by line and divides at ',' resulting in key (miner) and value (mapper) fields (size). The size of each miner is aggregated and yielded in the Combiner/Reducer phase.

Command: `python partc.py -r hadoop --output-dir part_c_a --no-cat-output hdfs://andromeda.eecs.qmul.ac.uk/data/ethereum/transactions/part-*.csv`

```
jmp01@itl211 ~/ecs765/ass1/partC> cat part_c.py
from mrjob.job import MRJob
import time

class partc(MRJob):

    def mapper(self,_,line):
        try:
            fields = line.split(',')
            if len(fields) == 9:
                size = float(fields[4])
                miner = fields[2]
                yield (fields[2], size)
        except:
            pass

    def combiner(self,key,val):
        yield (key,sum(val))

    def reducer(self,key,val):
        yield (key,sum(val))

if __name__ == '__main__':
    partc.run()

# python partc.py -r hadoop --output-dir partc --no-cat-output hdfs://andromeda.eecs.qmul.ac.uk/data/ethereum/transactions/part-*.csv
# job: http://andromeda.student.eecs.qmul.ac.uk:8088/proxy/application_1648683650522_7045/
```

hadoop fs -copyToLocal part_c_a inputfor10.txt

Top 10

If the input files are split along '\t' and have two fields, the mapper phase reads them line by line. We get value(None) and key(None) (miner, size). All pairs are sorted in descending order in the reducer phase using the sort function. We then repeat through the sorted values 10 times, yielding them, before stopping.

Command: python part_c_top_10.py inputfor10.txt > part_c_top10.txt

```
jmp01@itl211 ~/ecs765/ass1/partC> cat part_c_top10.py
from mrjob.job import MRJob

class top10_miner(MRJob):

    def mapper(self, _,line):
        try:
            fields = line.split('\t')
            if len(fields) == 2:
                address = fields[0]
                aggregate = float(fields[1])
                yield (None,(address,aggregate))
        except:
            pass

    def combiner(self, _, values):
        sorted_values = sorted(values, reverse = True, key = lambda x:x[1])
        for idx,value in enumerate(sorted_values):
            yield ("top", value)
            if idx >= 10:
                break

    def reducer(self, _, values):

        sorted_values = sorted(values, reverse = True, key = lambda x:x[1])
        for idx,value in enumerate(sorted_values):
            yield ("{} - {}".format(value[0],value[1]),None)
            if idx >= 10:
                break

if __name__ == '__main__':
    top10_miner.run()

# python part_c_top_10.py inputfor10.txt > part_c_top10.txt
```

Output:

```
jmp01@itl1211 ~/ecs765/ass1> cat part_c_top10.txt
"\0xea674fdde714fd979de3edf0f56aa9716b898ec8\" - 23989401188.0" null
"\0x829bd824b016326a401d083b33d09229333a830\" - 15010222714.0" null
"\0x5a0b54d5dc17e0aad383d2db43b0a0d3e029c4c\" - 13978859941.0" null
"\0x52bc44d5378309ee2abf1539bf71de1b7d7be3b5\" - 10998145387.0" null
"\0xb2930b35844a230f00e51431acae96fe543a0347\" - 7842595276.0" null
"\0x2a65aca4d5fc5b5c859090a6c34d164135398226\" - 3628875680.0" null
"\0x4bb96091ee9d802ed039c4d1a5f6216f90f81b01\" - 1221833144.0" null
"\0xf3b9d2c81f2b24b0fa0acaa865b7d9ced5fc2fb\" - 1152472379.0" null
"\0x1e9939daaad6924ad004c2560e90804164900341\" - 1080301927.0" null
"\0x61c808d82a3ac53231750dad13c77b59310bd9\" - 692942577.0" null
"\0xbcdfc35b86bedf72f0cda046a3c16829a2ef41d1\" - 610498611.0" null
jmp01@itl1211 ~/ecs765/ass1>
```

PART D. DATA EXPLORATION

Scam Analysis

The transaction dataset and the scams.json file were utilised as input files. Because we're dealing with 'json,' we've also imported the json library. Only if the input data set has fields with a length of 7 do we proceed (transactions data set). We'll use address as the key and value as the value from the transaction dataset. Similarly, we take address from the scams data set obtained from the for loop, assign value to it, and set 1 to differentiate it from transactions. The initial reduction phase determines whether the values come from transactions or json files. The reducer then sends the key value pair to the mapper, which includes categories and total values. After that, the second reduction step generates cumulative counts for each category.

Command: `python scam1.py -r hadoop --output-dir scam_part_1 --no-cat-output`

`hdfs://andromeda.eecs.qmul.ac.uk/data/ethereum/transactions/`

`hdfs://andromeda.eecs.qmul.ac.uk/data/ethereum/scams.json/`

```
jmp01@itl1211 ~/ecs765/ass1/partD> cat scam1.py
from mrjob.job import MRJob
from mrjob.step import MRStep
import json

class scams(MRJob):
    def mapper1(self, _, lines):
        try:
            fields = lines.split(",")
            if len(fields) == 7:
                address1 = fields[2]
                value = float(fields[3])
                yield address1, (value,0)

            else:
                line = json.loads(lines)
                keys = line["result"]

                for i in keys:
                    record = line["result"][i]
                    category = record["category"]
                    addresses = record["addresses"]

                    for j in addresses:
                        yield j, (category,1)

        except:
            pass

    def reducer1(self, key, values):
        tvalue=0
        category=None

        for k in values:
            if k[1] == 0:
                tvalue = tvalue + k[0]
            else:
                category = k[0]

        if category is not None:
            yield category, tvalue

    def mapper2(self,key,value):
        yield(key,value)
    def reducer2(self, key, value):
        yield(key,sum(value))

    def steps(self):
        return [MRStep(mapper = self.mapper1, reducer=self.reducer1), MRStep(mapper = self.mapper2,

if __name__ == '__main__':
    scams.run()
```


Output1:

```
jmp01@itl210 ~/ecs765/ass1/partD> cat scamout1.txt
"Scamming"      3.833616286244431e+22
"Fake ICO"      1.35645756688963e+21
"Phishing"      2.699937579408742e+22
"Scam" 0
. ....
```

Part 2

Similarly to the previous section, we import the json library and use the input files 'transactions.json' and 'scams.json'. To ensure that we are utilising the correct dataset, we check the length of the fields. The first mapper works in the same way as the previous code.

In the first reducer, we utilise address as a key, which is retrieved using a for loop, and category and status of the scam as the value, which is used to distinguish between them. Furthermore, it checks for the values and adds them if they are 1, else, the reducer will add them to categories and status as they are from the json file containing scam categories and values. The second mapper takes the values from the previous reduction and passes them on to the next reducer, which sums them up to tally the overall number of scams in a category.

Command: python scam2.py -r hadoop --output-dir scam_part_2 --no-cat-output

hdfs://andromeda.eecs.qmul.ac.uk/data/ethereum/transactions/

hdfs://andromeda.eecs.qmul.ac.uk/data/ethereum/scams.json/

```
jmp01@itl210 ~/ecs765/ass1/partD> cat scam2.py
from mrjob.job import MRJob
from mrjob.step import MRStep
import json

class scams2(MRJob):
    def mapper1(self, _, lines):
        try:
            fields = lines.split(",")
            if len(fields) == 7:
                address1 = fields[2]

                yield address1, (1,0)

        except:
            pass

        else:
            line = json.loads(lines)
            keys = line["result"]

            for i in keys:
                record = line["result"][i]
                category = record["category"]
                addresses = record["addresses"]
                status = record["status"]

                for j in addresses:
                    yield j, (2, category,status)

    def reducer1(self, key, values):
        tvalue=0
        category=None
        status = None
        for k in values:
            if k[0] == 1:
                tvalue = tvalue + k[0]
            else:
                category = k[1]
                status = k[2]
        if category is not None and status is not None:
            yield (status,category), tvalue

    def mapper2(self, key,value):
        yield(key,value)
    def reducer2(self, key, value):
        yield(key,sum(value))

    def steps(self):
        return [MRStep(mapper = self.mapper1, reducer=self.reducer1), MRStep(mapper = self.mapper2, reducer = self.reducer2)]

if __name__ == '__main__':
    scams2.run()
```

The output below shows the total number of frauds in various categories.

Output2:

```
[jmp01@itl210 ~/ecs765/ass1/partD> cat scamout2.txt
["Active", "Scamming"] 88444
["Inactive", "Phishing"] 22
["Offline", "Fake ICO"] 121
["Offline", "Phishing"] 7022
["Offline", "Scam"] 0
["Suspended", "Phishing"] 11
["Active", "Phishing"] 1584
["Offline", "Scamming"] 24692
["Suspended", "Scamming"] 56
jmp01@itl210 ~/ecs765/ass1/partD>
```

Contract types (identify contract types):

Extraction of features Top 5 transaction datasets with the most transactions

We can give award points to addresses that have the highest amount of transactions, and we can refer to them as loyal users.

2 jobs at the same time:

1. Here, the "address" field is yielded, and the reducer aggregates it to the number of transactions. The output is saved in the " Part_D_contracts_part_1" file.
2. In the second task, we used the sorted method in Python to apply the sorting algorithm to find the top 5 transactions.

Command: python Part_D_contracts_part_1.py -r hadoop --output-dir

Part_D_contracts_part_1 --no-cat-output

hdfs://andromeda.eecs.qmul.ac.uk/data/ethereum/transactions/

```
[jmp01@itl210 ~/ecs765/ass1/partD> cat Part_D_contracts_part_1.py
from mrjob.job import MRJob

class number_of_transactions(MRJob):
    def mapper(self, _, line):
        try:
            fields = line.split(",")
            if len(fields) == 7:
                key = fields[2]
                yield (key, 1)
        except:
            pass

    def combiner(self, k, l):
        yield k, sum(l)

    def reducer(self, u, t):
        yield (u, sum(t))

if __name__ == '__main__':
    number_of_transactions.run()

jmp01@itl210 ~/ecs765/ass1/partD>
```

Command: python Part_D_contracts_part_2.py -r hadoop --output-dir
Part_D_contracts_part_2 --no-cat-output
hdfs://andromeda.eecs.qmul.ac.uk/user/jmp01/Part_D_contracts_part_1/

```
[jmp01@it1210 ~/ecs765/ass1/partD] cat Part_D_contracts_part_2.py
from mrjob.job import MRJob

class number_of_transactions(MRJob):
    def mapper(self):
        try:
            fields = line.split()
            key = fields[0]
            value = int(fields[1])
            yield (None, (key, value))

        except:
            pass

    def reducer(self, _, val):
        sorted_values = sorted(val, reverse = True, key = lambda tup: tup[1])
        j = 0
        for i in sorted_values:
            yield i
            j += 1
            if j >= 5:
                break

if __name__ == '__main__':
    number_of_transactions.run()
[jmp01@it1210 ~/ecs765/ass1/partD]
```

Output:

```
"\"0x8d12a197cb00d4747a1fe03395095ce2a5cc6819\" 11089018
\"0x2a0c0dbec7e4d658f48e01e3fa353f44050c208\" 7160876
\"0x3f5ce5fbfe3e9af3971dd833d26ba9b5c936f0be\" 5952175
\"0x174bfa6600bf90c885c7c01c7031389ed1461ab9\" 5797049
\"0x06012c8cf97bead5deae237070f9587f8e7a266d\" 3995252
```

Is_erc721_contract:

Erc721 is an ethereum token that may be used to identify contracts that adhere to the erc721 contract rules. Smart contracts are another name for them. By simply adding a "if" condition that checks for the truth value in the relevant column of the dataset, we can filter the contract by determining if it is an er721 contract.

Command: python Part_D_Is_erc721_contract.py -r hadoop --output-dir
Part_D_Is_erc721_contract --no-cat-output
hdfs://andromeda.eecs.qmul.ac.uk/data/ethereum/contracts/

```
[jmp01@itl210 ~/ecs765/ass1/partD] cat Part_D_contracts_part_2.py
from mrjob.job import MRJob

class number_of_transactions(MRJob):
    def mapper(self):
        try:
            fields = line.split()
            key = fields[0]
            value = int(fields[1])
            yield (None, (key, value))

        except:
            pass

    def reducer(self, _, val):
        sorted_values = sorted(val, reverse = True, key = lambda tup: tup[1])
        j = 0
        for i in sorted_values:
            yield i
            j += 1
            if j >= 5:
                break

if __name__ == '__main__':
    number_of_transactions.run()

[jmp01@itl210 ~/ecs765/ass1/partD] nano Part_D_Is_erc721_contract.py
[jmp01@itl210 ~/ecs765/ass1/partD] cat Part_D_Is_erc721_contract.py
from mrjob.job import MRJob

class erc721_contract(MRJob):
    def mapper(self, _, line):
        try:
            fields = line.split(',')
            if len(fields) == 5:
                if fields[2] == "true":
                    yield fields[0], 1

        except:
            pass

    def reducer(self, u, t):
        yield (u, sum(t))

if __name__ == '__main__':
    erc721_contract.run()

[jmp01@itl210 ~/ecs765/ass1/partD] █
```

Output:

```
0x000edf42475e7ceb32f82ee11f9733231a67b2be" 1
0x00136a57574c805312d4ee875b0bc2e56984f00d" 1
0x003203f31def0aa63b895b8599c9b81ddce8939b" 1
0x003ad9c18bc279f40632e7e5de2fd213931215d0" 1
0x0063f8d3537ec9cd23b08357494d3e0ee63a8f4a" 1
0x00650ea64d5c226755a6a6976a774a6f2fdf8c13" 1
0x006bf2f5cc930eed50e14c66605ac95767133bec" 1
0x00b218fd2db515392f2a2d78b658c913f2eebf2d" 1
0x00bf70e1ddfb8984d0af9af4b29ad3ec40d4b84e" 1
0x00d2be2c7e509515cdba64051a643a5220f1f241" 1
0x0111ac7e9425c891f935c4ce54cf16db7c14b7db" 1
0x01144cad63687f9d06c46854c810366b28a84d73" 1
0x0131c575be43586c5346435273035c8709a3fd34" 1
0x01d8a618ba50232e73b7e0c3254bfa0802a75b56" 1
0x01f96d3376acae2fc4f9077f03fa9d997797ed5c" 1
0x0211621fb78c85ee44b12d2f04e2e2f16334a9fe" 1
0x0213c2c85315299c6074fc4a6d779121284e9f42" 1
.....
```

Outflow addresses:

Simply apply wordcount to the transaction dataset to obtain a unique outflow address. This allows us to retrieve all of the transaction dataset's unique addresses.

Command: `python Part_D_outflow.py -r hadoop --output-dir Part_D_outflow --no-cat-output hdfs://andromeda.eecs.qmul.ac.uk/data/ethereum/transactions/`

```
[jmp01@itl210 ~/ecs765/ass1/partD> cat Part_D_outflow.py
from mrjob.job import MRJob

class contract(MRJob):
    def mapper(self, _, line):
        try:
            fields = line.split(",")
            if len(fields) == 7:
                yield fields[2], 1

        except:
            pass

    def reducer(self, k, v):
        yield k, sum(v)

if __name__ == '__main__':
    contract.run()

jmp01@itl210 ~/ecs765/ass1/partD> █
```

Output:

```
"0x014d3c3de0696c4e1662cd5e913d833e9181af89" 1
"0x014d3c62d7d8d41936773b186be11af.2d2e6ac16" 1
"0x014d3cd32875eb5fab7cdc383398fe8b8.28181d8" 3
"0x014d3cd43567dac810aba9e0eec6fff084bb0159d" 1
"0x014d3d3cd0d2f6334a7d70b26a72038118875df9" 3
"0x014d3d762587bc8f33fd2eff618cc8505d45a290" 5
"0x014d3dc71911e397602c3991b830bdal7c5a8a71" 1
"0x014d3de457357bcf065d5897e905cac292bc452c" 2
"0x014d3e9e5216d77067af40fce741e0aaf2bedc90" 1
"0x014d3eb929c0f449345c87dcde4a4e13f91c7246" 1
"0x014d3ed7a70692491a832aac8.2d.241ed967c8401" 1
"0x014d3f101261d1b7221492946123017d786a5f6d" 1
"0x014d3f2be1c647a725742cdcb917f235bcd0e2074" 64
"0x014d405b8dcd1723eec901b078d481928805d103" 15
"0x014d40d44c31f7f4f1795bf4eb0ca385cdfa92d" 1
"0x014d40f86e343c298b03446a8d86f7022af1465e" 1
"0x014d42589bfbde32b84190d7fdbce618c4a4d9c2" 1
"0x014d42bfcbe9c00fb956af2b837caf231ffebd8d" 1
"0x014d43a3b18f0944fee6622a5815ba0321d6440a" 9
"0x014d4517ecd1e4bb9e7c0c7020de16b9f68c7faa" 1
"0x014d45d33b131b4ecda18c3a4924a7ac15c58220" 1
"0x014d46c6eff0e04a2.278f017671f73ea7ed2ca9d" 1
"0x014d47e551cac04db2ab4980d81f29cb38ac16b3" 1
"0x014d47fc325050fe37347d96e3a5726284ef664a" 1
"0x014d48483f5b111e397974a75322dbbd1a6b0cce" 1
"0x014d495cd68302745445fda7cd42ce1d9e033f29" 1
"0x014d49efeb2a7fbca43705b4eaf957a02813df19" 2
"0x014d4a2cd11519245ca78d30b68a5e935e518e12" 22
"0x014d4a2fed1aabe08e82edfa28343c94d3800442" 1
"0x014d4b916c71cdcb0122d8c5605ff7a1d500a518" 1
"0x014d4be25067f7dde1187b5dffa9c2ace3e395d6" 81
"0x014d4bf25039f4f47b74ea73db42399b97288ca2" 2
"0x014d4bfa1886b969ede482fa7af5045f333dd59b" 2
"0x014d4cf0721be83114alac9d321664cc3084cd3c" 2
```

MISCELLANEOUS

Gas Guzzlers:

Using MapReduce, obtain the average Gas Price and Average Gas Limit for each month in the dataset by running the following two tasks.

GasPrice.py

```
[jmp01@itl210 ~/ecs765/ass1/partD> cat GasPrice.py
from mrjob.job import MRJob
import time

class Gas(MRJob):

    def mapper(self,_,line):
        try:
            fields = line.split(',')
            val = float(fields[5])
            date = time.localtime(float(fields[6]))
            if len(fields) == 7:
                yield ((date.tm_mon,date.tm_year),(1,val))    #Number of transaction

        except:
            pass

    def combiner(self,key,val):
        count = 0
        total = 0
        for v in val:
            count+=v[0]
            total+=v[1]
        yield (key,(count,total))

    def reducer(self,key,val):
        count = 0
        total = 0
        for v in val:
            count+=v[0]
            total+=v[1]
        yield (key,(total/count))

if __name__=='__main__':
    Gas.run()
[jmp01@itl210 ~/ecs765/ass1/partD> █
```

GasLimit.py

```
[jmp01@itl210 ~/ecs765/ass1/partD> cat GasLimit.py
from mrjob.job import MRJob
import time

class Gas_limit(MRJob):

    def mapper(self,_,line):
        try:
            fields = line.split(',')
            val = float(fields[4])
            date = time.localtime(float(fields[6]))
            if len(fields) == 7:
                yield ((date.tm_mon,date.tm_year),(1,val))    #Number of transaction

        except:
            pass

    def combiner(self,key,val):
        count = 0
        total = 0
        for v in val:
            count+=v[0]
            total+=v[1]
        yield (key,(count,total))

    def reducer(self,key,val):
        count = 0
        total = 0
        for v in val:
            count+=v[0]
            total+=v[1]
        yield (key,(total/count))

if __name__=='__main__':
    Gas_limit.run()

[jmp01@itl210 ~/ecs765/ass1/partD> █
```

Output:

```
jmp01@itl210 ~/ecs765/ass1/partD> cat Gas_Price_output.txt
[1, 2017] 22507570807.719795
[1, 2019] 14611816445.785261
[10, 2015] 53898497955.07804
[10, 2017] 17509171844.77064
[11, 2016] 24634294365.279037
[11, 2018] 16034859008.681648
[12, 2015] 55899526672.35498
[12, 2017] 33423472930.407898
[2, 2016] 69180681134.38954
[2, 2018] 23636574203.828873
[3, 2017] 23232083087.910202
[3, 2019] 18091340267.2465
[4, 2016] 23359978331.676765
[4, 2018] 13149523170.45877
[5, 2017] 23568661138.035046
[5, 2019] 14480574461.419468
[6, 2016] 23021831286.57488
[6, 2018] 16536952425.540333
[7, 2017] 25463022668.143833
[8, 2016] 22407628763.365406
[8, 2018] 18478650928.737164
[9, 2015] 56512934320.01927
[9, 2017] 30676555381.728436
[1, 2016] 56596270931.316185
[1, 2018] 52106060636.844185
[10, 2016] 32113146198.891758
[10, 2018] 14527572489.59486
[11, 2015] 53607614201.79755
[11, 2017] 15312465314.693539
[12, 2016] 50318068074.687996
[12, 2018] 16338844844.014513
[2, 2017] 23047230327.254387
[2, 2019] 28940599438.1487
[3, 2016] 32805967466.947445
[3, 2018] 15554999714.874079
[4, 2017] 22357075153.737774
[4, 2019] 11573133401.384796
[5, 2016] 23747073761.79113
[5, 2018] 17414613148.43696
[6, 2017] 30201664896.657127
[6, 2019] 15067557451.33386
[7, 2016] 22619213302.825947
[7, 2018] 27520561081.0211
[8, 2015] 160356354969.0592
[8, 2017] 25903650367.89669
[9, 2016] 25262249340.11566
[9, 2018] 15208159827.250359


jmp01@itl210 ~/ecs765/ass1/partD> █
```

```
jmp01@itl210 ~/ecs765/ass1/partD> cat Gas_Limit_output.txt
[1, 2017] 144585.8323444452
[1, 2019] 213715.9605048934
[10, 2015] 124898.76345065859
[10, 2017] 139870.05913735816
[11, 2016] 118474.20698056064
[11, 2018] 224208.95181702235
[12, 2015] 202883.632316504
[12, 2017] 125269.51047898784
[2, 2016] 130175.52141758327
[2, 2018] 131824.6970378884
[3, 2017] 147228.19778654756
[3, 2019] 248217.42927149552
[4, 2016] 95083.47755021007
[4, 2018] 236204.97904008476
[5, 2017] 152449.08294623252
[5, 2019] 183094.5693197547
[6, 2016] 130112.21626908527
[6, 2018] 180795.02085380582
[7, 2017] 132790.34469405047
[8, 2016] 119958.22676283357
[8, 2018] 185752.46188363043
[9, 2015] 99880.37666969116
[9, 2017] 146371.34922749794
[1, 2016] 140463.7812759377
[1, 2018] 102167.12354317823
[10, 2016] 117078.86856704934
[10, 2018] 202550.24404083
[11, 2015] 209520.94340378218
[11, 2017] 146613.7212948582
[12, 2016] 134463.8302327048
[12, 2018] 228197.0665342265
[2, 2017] 185771.26611789103
[2, 2019] 246776.93062665823
[3, 2016] 103585.96288806947
[3, 2018] 156114.3435117928
[4, 2017] 156866.61651822302
[4, 2019] 210658.60478696763
[5, 2016] 119737.9383389679
[5, 2018] 149300.75953211795
[6, 2017] 129488.39504908556
[6, 2019] 153578.42011967467
[7, 2016] 114490.23118436741
[7, 2018] 160946.32119566188
[8, 2015] 73133.53050098693
[8, 2017] 139510.6310349726
[9, 2016] 139595.43180709172
[9, 2018] 202264.51543298518

jmp01@itl210 ~/ecs765/ass1/partD> █
```


Bar charts & Observation

Gas_Price vs Time :



According to the graph above, the average price of gas has reduced from 2015 to 2019, and we can also see that each year, prices have risen in the last few months or the first few months of the year.

Contract_gas vs Time :



Observation from the graph above - Despite some ups and downs, there has been a general pattern of steadily increasing gas requirements (from 80000 to 250000 approx)